

Complete guide to Association Rules (2/2)

Algorithms that help you shop faster and smarter



Anisha Garg · [Follow](#)

Published in Towards Data Science

8 min read · Sep 17, 2018



Listen



Share



In this blog, I will discuss the algorithms that enable efficient extraction of association rules from a list of transactions. Part 1 of this blog covers the terminology and concepts that form the foundation of association rule mining. Motivation behind this whole concept and meaning of some basic terms is explained there. Here are very brief definitions of some terms from the previous part.

1. Association Rule: Ex. $\{X \rightarrow Y\}$ is a representation of finding Y on the basket which has X on it
2. Itemset: Ex. $\{X, Y\}$ is a representation of the list of all items which form the association rule
3. Support: Fraction of transactions containing the itemset

4. Confidence: Probability of occurrence of $\{Y\}$ given $\{X\}$ is present
5. Lift: Ratio of *confidence* to baseline probability of occurrence of $\{Y\}$


Now that we are familiar with these terms, let's proceed ahead with extracting the rules from a list of transactions. The very first step in the process is to get a list of all the items occurring at least once in our set of transactions.

The challenge is the mining of important rules from a massive number of association rules that can be derived from a list of items.

Remember, rule-generation is a two step process. First is to generate an itemset like $\{\text{Bread, Egg, Milk}\}$ and second is to generate a rule from each itemset like $\{\text{Bread} \rightarrow \text{Egg, Milk}\}$, $\{\text{Bread, Egg} \rightarrow \text{Milk}\}$ etc. Both the steps are discussed below.

1. Generating itemsets from a list of items

First step in generation of association rules is to get all the *frequent* itemsets on which binary partitions can be performed to get the antecedent and the consequent. For example, if there are 6 items $\{\text{Bread, Butter, Egg, Milk, Notebook, Toothbrush}\}$ on all the transactions combined, itemsets will look like $\{\text{Bread}\}$, $\{\text{Butter}\}$, $\{\text{Bread, Notebook}\}$, $\{\text{Milk, Toothbrush}\}$, $\{\text{Milk, Egg, Vegetables}\}$ etc. Size of an itemset can vary from one to the total number of items that we have. Now, we seek only *frequent* itemsets from this and not all so as to put a check on the number of total itemsets generated.



Frequent itemsets are the ones which occur at least a minimum number of times in the transactions. Technically, these are the itemsets for which *support* value (fraction of transactions containing the itemset) is above a minimum threshold — *minsup*.

So, {Bread, Notebook} might not be a frequent itemset if it occurs only 2 times out of 100 transactions and $(2/100) = 0.02$ falls below the value of minsup.

A brute force approach to find frequent itemsets is to form all possible itemsets and check the support value of each of these. *Apriori principle* helps in making this search efficient. It states that

All subsets of a frequent itemset must also be frequent.

This is equivalent to saying that number of transactions containing items {Bread, Egg} is greater than or equal to number of transactions containing {Bread, Egg, Vegetables}. If the latter occurs in 30 transactions, former is occurring in all 30 of them and possibly will occur in even some more transactions. So if support value of {Bread, Egg, Vegetables} i.e. $(30/100) = 0.3$ is above minsup, then we can be assured that support value of {Bread, Egg} i.e. $(>30/100) = >0.3$ is above minsup too. This is called the **anti-monotone property of support** where if we drop out an item from an itemset, support value of new itemset generated will either be the same or will increase.

Apriori principle is a result of anti-monotone property of support.

Apriori principle allows us to prune all the supersets of an itemset which does not satisfy the minimum threshold condition for support. For example, if {Milk, Notebook} does not satisfy our threshold of minsup, an itemset with any item added to this will never cross the threshold too.

The methodology that results is called the *apriori algorithm*. Steps involved are:

Ref: <https://annalyzin.files.wordpress.com/2016/04/association-rules-apriori-tutorial-explanation.gif>

Generate all frequent itemsets (support \geq minsup) having only one item. Next, generate itemsets of length 2 as all possible combinations of above itemsets. Then, prune the ones for which support value fell below minsup.

Now generate itemsets of length 3 as all possible combinations of length 2 itemsets (that remained after pruning) and perform the same check on support value.

We keep increasing the length of itemsets by one like this and check for the threshold at each step.

As can be seen from the graphic, pruning of infrequent itemsets reduces the number of itemsets to be considered by more than half! This proportion of reduction in computational power becomes more and more significant as the number of items increases.

This proportion also depends on the minimum support threshold (minsup) that we pick up which is completely subjective to the problem at hand and can be based on past experience.

2. Generating all possible rules from the frequent itemsets

Once the frequent itemsets are generated, identifying rules out of them is comparatively less taxing. Rules are formed by binary partition of each itemset. If {Bread,Egg,Milk,Butter} is the frequent itemset, candidate rules will look like:

(Egg, Milk, Butter \rightarrow Bread), (Bread, Milk, Butter \rightarrow Egg), (Bread, Egg \rightarrow Milk, Butter), (Egg, Milk \rightarrow Bread, Butter), (Butter \rightarrow Bread, Egg, Milk)

From a list of all possible candidate rules, we aim to identify rules that fall above a minimum confidence level (*minconf*). Just like the anti-monotone property of support, *confidence of rules generated from the same itemset also follows an anti-monotone property*. It is anti-monotone with respect to the number of elements in consequent.

This means that confidence of $(A,B,C \rightarrow D) \geq (B,C \rightarrow A,D) \geq (C \rightarrow A,B,D)$. To remind, confidence for $\{X \rightarrow Y\} = \text{support of } \{X,Y\} / \text{support of } \{X\}$

As we know that support of all the rules generated from same itemset remains the same and difference occurs only in the denominator calculation of confidence. As number of items in X decrease, support{X} increases (as follows from the anti-monotone property of support) and hence the confidence value decreases.

An intuitive explanation for the above will be as follows. Consider F1 and F2:

F1 = fraction of transactions having (butter) also having (egg, milk, bread)

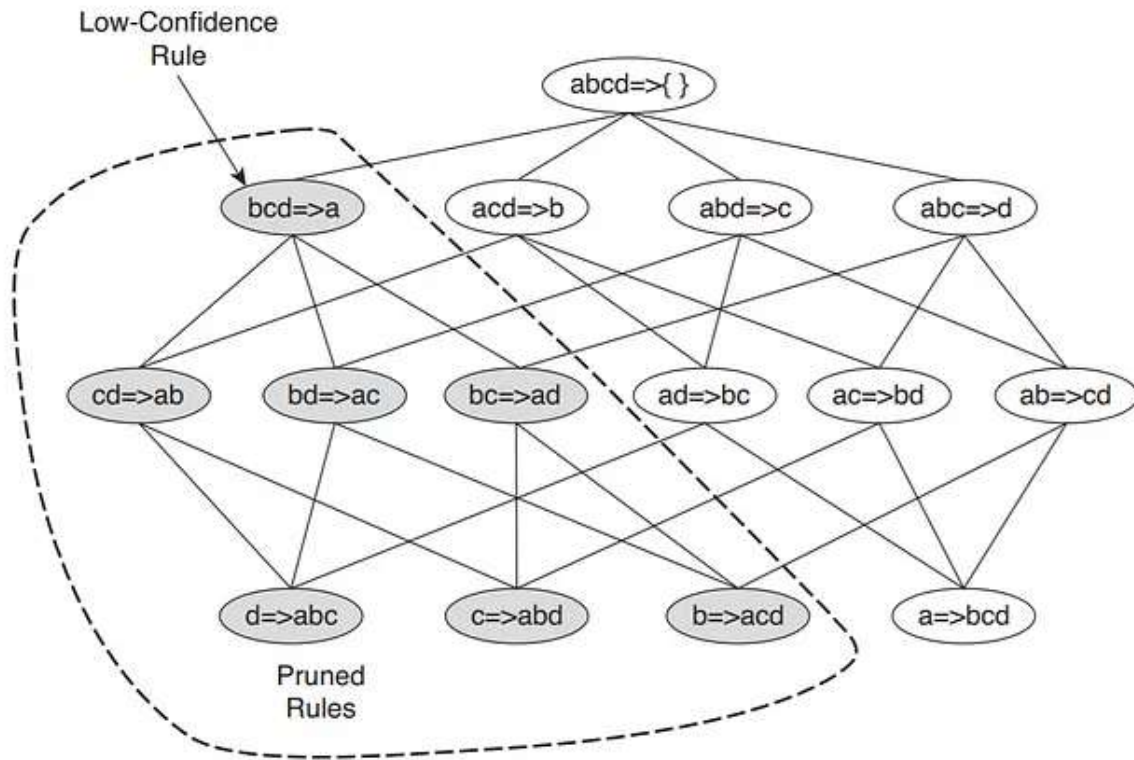
There will be many transactions having butter and all three of egg, milk and bread will be able to find place only in a small number of those.

F2 = fraction of transactions having (milk, butter, bread) also having (egg)

There will only be a handful of transactions having all three of milk, butter and bread (as compared to having just butter) and there will be high chances of having egg on those.

So it will be observed that $F1 < F2$. Using this property of confidence, pruning is done in a similar way as was done while looking for frequent itemsets. It is illustrated in the figure below.





Ref: <https://www-users.cs.umn.edu/~kumar001/dmbook/ch6.pdf>


We start with a frequent itemset $\{a,b,c,d\}$ and start forming rules with just one consequent. Remove the rules failing to satisfy the minconf condition. Now, start forming rules using a combination of consequents from the remaining ones. Keep repeating until only one item is left on antecedent. This process has to be done for all frequent itemsets.

Here again, minimum confidence threshold that we pick up is completely subjective to the problem at hand.

With these two steps, we have identified a set of association rules which satisfy both the minimum support and minimum confidence condition. The number of such rules obtained will vary with the values of *minsup* and *minconf*. Now, this subset of rules thus generated can be searched for highest values of lift to make business decisions. There are some nicely built libraries in R to fetch association rules from transactions by putting in *minsup* and *minconf* as parameters. They also provide capabilities to visualize the same. This article [here](#) discusses the whole process step-by-step with the code.

Before concluding, I would introduce two more terms, maximal frequent itemset and closed frequent itemset which are used as a compact representation of all the frequent itemsets.

Maximal frequent itemset: *It is a frequent itemset for which none of the immediate supersets are frequent.* This is like a frequent itemset X to which no item y can be added such that {X,y} still remains above minsup threshold.



Different cases for itemset {Bread, Butter}

Closed frequent itemset: *It is a frequent itemset for which there exists no superset which has the same support as the itemset.* Consider an itemset X. If ALL occurrences of X are accompanied by occurrence of Y, then X is NOT a closed set. (Refer the figure below for example)

Maximal frequent itemsets are valuable because they are the most compact form of representation of frequent itemsets.

All the frequent itemsets can be derived as the subsets of maximal frequent itemsets. However, information on support of the subsets is lost. If this value is required, closed frequent itemset is another way to represent all the frequent itemsets.

Closed itemsets help in removing some redundant itemsets while not losing information about the support values.

Calculation of support values of non-closed itemsets from closed itemsets is another algorithm, details of which are out of scope of this article.

I have tried to cover all the important terms and concepts related to mining of association rules through this blog going into details wherever necessary. Following are one line summaries for few terms introduced in this process —

1. Association rule mining: (a) Itemset generation, (b) Rule generation
2. Apriori principle: All subsets of a frequent itemset must also be frequent
3. Apriori algorithm: Pruning to efficiently get all the frequent itemsets
4. Maximal frequent itemset: none of the immediate supersets are frequent
5. Closed frequent itemset: none of the immediate supersets have the same value of support

I hope you enjoyed reading this and have more clarity in thoughts than before. If you are interested in reading more details, refer [this](#) resource. Let me know of your thoughts/questions through the comments.

Do read the introduction in first blog [here](#) if you haven't already!

[Data Science](#)[Machine Learning](#)[Concept](#)[Analytics](#)[Analysis](#)[Follow](#)

Written by Anisha Garg

605 Followers · Writer for Towards Data Science

Data Scientist | UT Austin, IIT Bombay alum

More from Anisha Garg and Towards Data Science



 Anisha Garg in Towards Data Science

Complete guide to Association Rules (1/2)

Algorithms that help you shop faster and smarter

7 min read · Sep 3, 2018

 2.1K

 16







Jacob Marks, Ph.D. in Towards Data Science

How I Turned My Company's Docs into a Searchable Database with OpenAI

And how you can do the same with your docs


15 min read · Apr 25



4K



50



Khuyen Tran in Towards Data Science

Stop Hard Coding in a Data Science Project—Use Config Files Instead

And How to Efficiently Interact with Config Files in Python



· 6 min read · May 26




1.5K



19





Anisha Garg in Towards Data Science

Basic statistics for every Data Scientist

How confident are you in your estimates?!

6 min read · Nov 11, 2019




391



See all from Anisha Garg

See all from Towards Data Science


Recommended from Medium



 Matt Crooks in Towards Data Science

A Review of Propensity Score Modelling Approaches

A review of different approaches to using propensity scores in causal inference modelling

 · 11 min read · May 17

 143  1



 Rahul S. (Aaweg I, ShabdAaweg)

Enhancing Machine Learning Projects: Strategies for Effective Data Handling and Model Performance


Machine learning has revolutionized numerous industries, from finance to healthcare, by enabling the development of intelligent systems...

★ · 10 min read · Jun 10

 22 



Lists



Predictive Modeling w/ Python

18 stories · 9 saves




Practical Guides to Machine Learning

10 stories · 22 saves



Natural Language Processing

345 stories · 7 saves



New_Reading_List

173 stories · 1 save



Santiago Rodrigues Manica in Epidence

Sample Size Matters

How to calculate sample size using R and Python

★ · 4 min read · Feb 5



10

 Data Overload

Market Basket Analysis: Techniques, Applications, and Benefits for Retailers

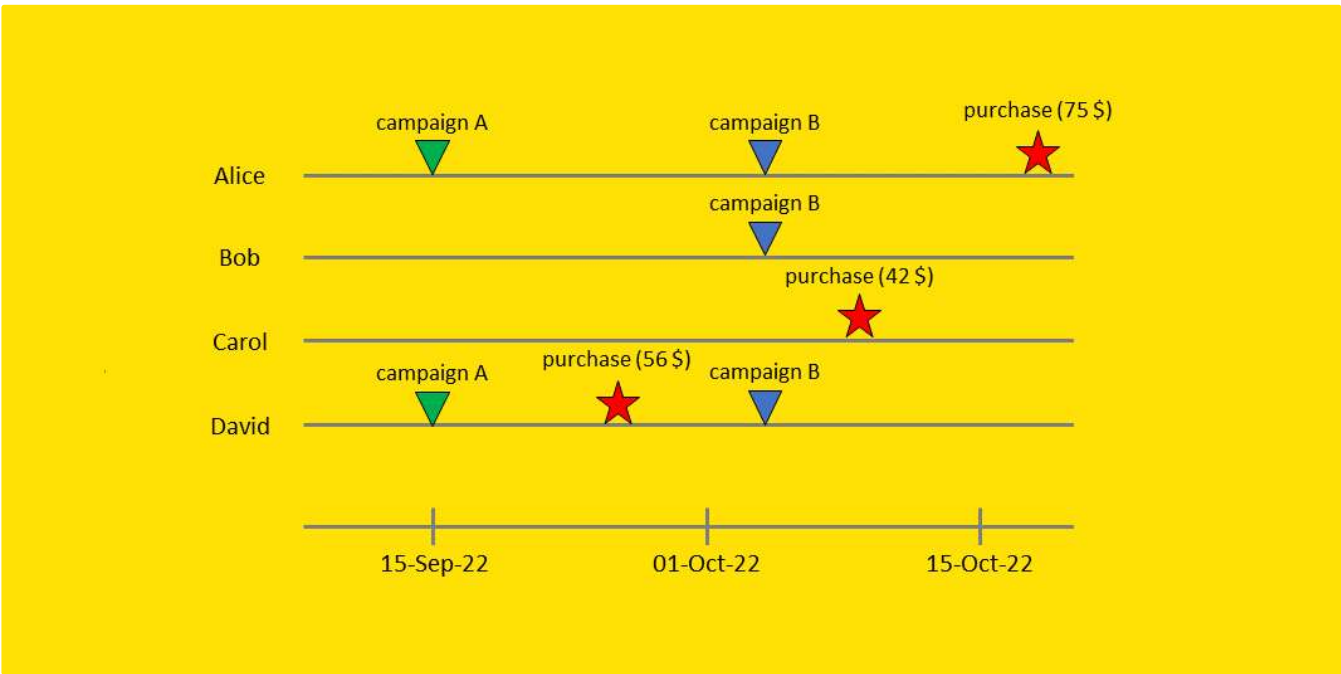
Market basket analysis is a data mining technique used by retailers to analyze the purchase behavior of customers. The technique identifies...


★ · 5 min read · Mar 9



1






 Samuele Mazzanti in Towards Data Science


Using Causal ML Instead of A/B Testing

In complex environments, Causal ML is a powerful tool because it is more flexible than A/B Testing, and it doesn't require strong...

★ · 9 min read · Nov 29, 2022

 937  15



 Leonie Monigatti in Towards Data Science

10 Exciting Project Ideas Using Large Language Models (LLMs) for Your Portfolio

Learn how to build apps and showcase your skills with large language models (LLMs). Get started today!

★

• 11 min read

• May 15

 1.4K

 9



See more recommendations